

C4 Integer Overflows & Targeted Overwrites

1. Security Issues with Integers

1.1 Background on Integers

Natural Numbers

- The natural numbers (counting numbers) \mathbb{N} are defined by the Peano postulates:
 - 1 is a member of the set \mathbb{N}
 - If n is a member of \mathbb{N} , then the **successor** $s(n)$ belongs to \mathbb{N}
 - 1 is not the successor of any element in \mathbb{N}
 - If $s(n) = s(m)$, then $n = m$
 - If $m \in \mathbb{N}$ is not 1, then there exists a $n \in \mathbb{N}$ such that $s(n) = m$
 - A subset of \mathbb{N} which contains 1, and which contains $n + 1$ whenever it contains n , must equal \mathbb{N}
- You can write $n + 1$ instead of $s(n)$
- $a, b \in \mathbb{N}$, if $b = 1$, then $a + b = s(a)$, else take $c \in \mathbb{N}$ so that $s(c) = b$ and define $a + b = s(a + c)$

1.2 Integer Overflows

- **Unsigned 8-bit integers**
 $255 + 1 = 0$ $16 \times 17 = 16$ $0 - 1 = 255$
- **Unsigned 64-bit integers**
 $2^{63} + 2^{63} = 0$
- **Signed 8-bit integers**
 $127 + 1 = -128$ $-128 \div -1 = -128$
- Can lead to buffer overflows (bof)
 - OS checks string lengths to defend against bof

- "safe" signed addition prevents addition of arguments with the same sign

1.3 Integer Mismatches

- Declare all integers as unsigned integers unless negative numbers are needed (compilers will issue warning if signed-unsigned mismatch occurs)
- **Truncation:** input UID as signed integer, check value $\neq 0$, truncate to unsigned short integer $0x10000 \rightarrow 0x0000$ (root!)
- Integer overflows happen when behaviours of abstraction and implementation diverge
 - Place checks in code or use libraries that detect such situation and flag error so that caller can handle these error messages correctly

2. Format String Attacks

- Overwriting arbitrary pointer with an arbitrary value
- Bof overwrites all position between buffer and target which may crash a program before returning
- In format string attacks, `printf()` does the attacker's job
- In double-free attacks, `malloc()` does the attacker's job

2.1 printf()

<code>printf(char *, ...);</code>	creates a formatted string and writes it to standard out I/O stream
<code>fprintf(FILE *, char *, ...);</code>	creates a formatted string and writes it to a libc FILE I/O stream
<code>sprintf(char *, char *, ...);</code>	creates a formatted string and writes it to a location in memory
<code>snprintf(char *, size_t, char *, ...);</code>	creates a formatted string and writes it to a location in memory, with a maximum string size

<code>%i,</code> <code>%d</code>	<code>int, short, char</code>	Integer value of argument in decimal notation
<code>%u</code>	<code>unsigned int, short, char</code>	Value of argument as unsigned integer in decimal notation
<code>%x</code>	<code>unsigned int, short, char</code>	Value of argument as unsigned integer in hexadecimal notation
<code>%s</code>	<code>char *, char[]</code>	Character string pointed to by the argument
<code>%p</code>	<code>(void *)</code>	Value of the pointer in hexadecimal notation
<code>%n</code>	<code>(int *)</code>	Number of bytes output so far, stored in an argument that is a pointer to an integer

```
char buf[10];
int ctr, x = 0;
snprintf(buf, sizeof buf,
          "%.100d%n", x, &ctr);
printf("counter: %d\n", ctr);
```

.precision: here 100 digits for rendering x

- **Counter** is incremented to 100, although only 10 characters are actually written to the buffer

- `snprintf()` saves number of characters printed out to address allocated to `ctr`
- `printf("%.100d%n", x, &ctr)` works the same way, although there is no maximum string size

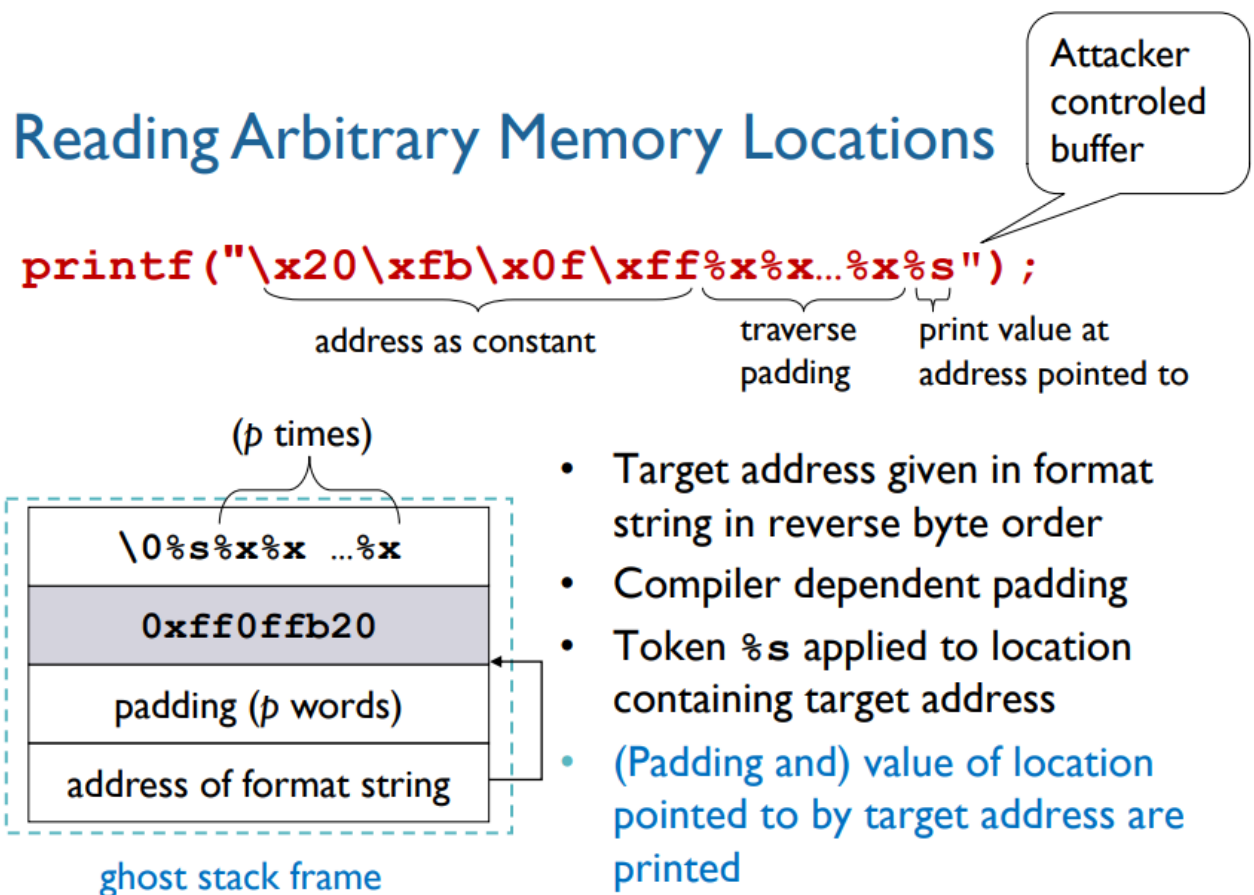
Execution

- Format string & arguments provided to `printf()` allocated on the stack

- Stack frame for `printf()` contains arguments passed to the function
- **Assume** address of format string passed in as argument
- No. of arguments calculated by counting **format tokens (e.g. %c, %d)** in the format string
 - If no. of tokens > no. of arguments, memory positions next in stack frame will be printed (in 64 bit linux, parameters passing is done in this order: RDI, RSI, RDX, RCX, R8, R9, remaining from the stack)

Reading Arbitrary Memory Locations

- Identify format string function (e.g. `printf()`) where you get to specify format string
- When `printf()` is called without format tokens, an attack can pass its own format tokens to create a *ghost* stack frame
- To read from a chosen memory address, put the address as a constant in the format string
- Construct format string so that this constant becomes argument for a `%s` (char array) format token in the ghost frame



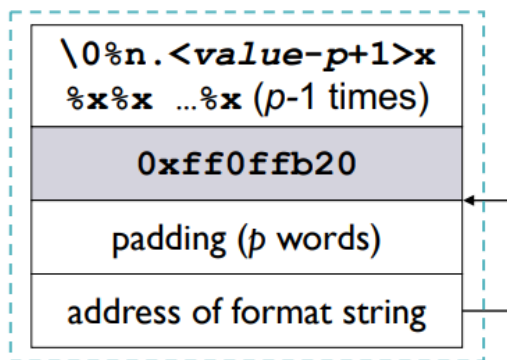
Writing to Arbitrary Memory Locations

- Include format tokens to traverse memory locations from location the first argument is expected to be at to the address of the format string

- Use `.precision` to set counter value of your choice

```
printf("\x20\xfb\x0f\xff%x%x...%x%.<value-p+1>x%n");
```

address as constant
traverse padding
increment counter
write value to address pointed to



ghost stack frame

- Target address given in format string in reverse byte order
- Compiler dependent padding
- `%.<value-p+1>x` increments counter to desired value
- Token `%n` applied to location containing target address
- **value** is written to the location pointed to by the target address

(similar vid: <https://www.youtube.com/watch?v=CyazDp-Kkr0>)

3. Heap Memory Allocation

3.1 Malloc() and Free()

Malloc()

- `void * malloc(size)` returns pointer to newly allocated block of `size` bytes
 - contents of block are not initialised
 - returns null pointer if call fails

Free()

- `void free (void *ptr)`
 - `ptr` must come from previously called `malloc, calloc, realloc`
 - no operation performed if `ptr` is null
 - if called again, behaviour is undefined
 - **not atomic (not done in one instruction call)**

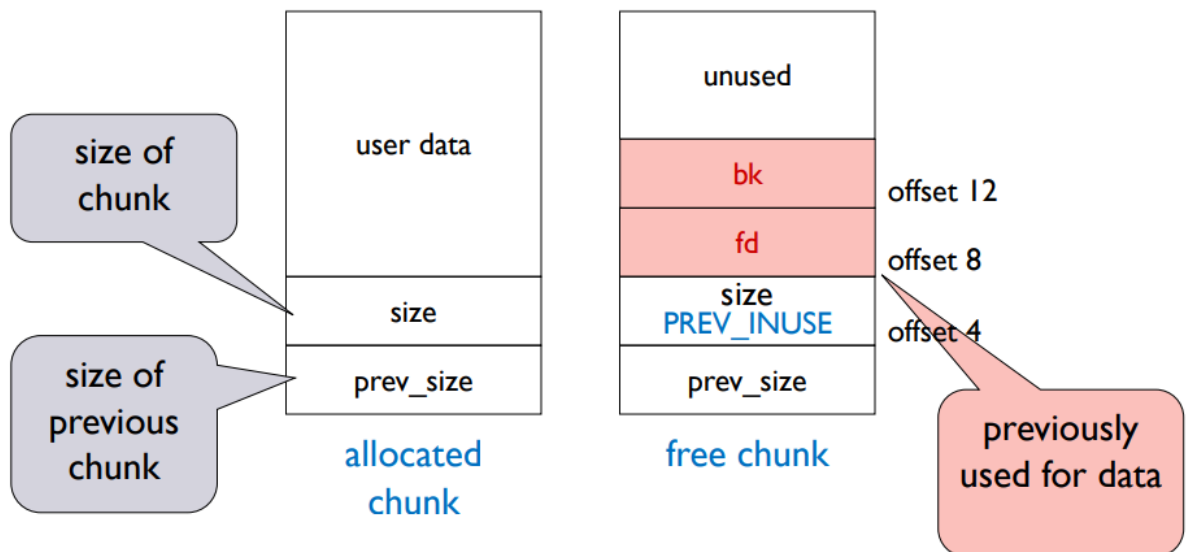
3.2 Doubly Linked Lists for Memory Management

Memory Organisation

- Memory is divided into chunks which contain user data and control data
 - Chunks allocated by malloc contain boundary tags
 - Free chunks are placed in bins (doubly linked list)
 - Tags of a free chunk contain a forward (fd) and backward (bk) pointer to its neighbours in the bin (can have allocated chunks in between)

Allocated and Free Chunks

```
struct chunk {  
    int prev_size;  
    int size;  
    /* used of if free */  
    struct chunk *fd;  
    struct chunk *bk;  
}
```



- Allocated chunks do not have pointers to neighbouring chunks
- Values for `size` is always given in multiples of 8 with the 3 LSBs used as control flags (`PREV_INUSE (0x1)`, `IS_MAPPED (0x2)`, some libraries use the 3rd bit)
 - `[size][??][mapped][prev_chunk_inuse]`
 - When a chunk is freed it is combined into a single chunk with its neighbours
 - `PREV_INUSE` is used to check if the previous chunk is in use, if not, combine both chunks

Two Views on Memory

- **“Topological” (physical) view:** location of chunks in memory
 - Chunks ordered by their addresses
 - We can talk about the **“chunk above”** and the **“chunk below”**
 - This view matters when chunks are coalesced
- **“Logical” view:** location of chunks in a bin
 - Chunks ordered by their position in a list
 - We can talk about **“previous chunk”** and **“next chunk”**
 - This view matters when chunks are allocated and freed

Bin Management

- Doubly-linked lists of free chunks **ordered by increasing size** to facilitate fast smallest-first search

- `unlink()` is used to allocate a buffer during `malloc()`
- `frontlink()` is used to insert a freed chunk to the right position (by size)

Frontlink()

- Simplified version

```
#define frontlink(A, P, S, IDX, BK, FD) ...
[1] FD = start_of_bin(IDX);
[2] while ( FD != BK && S < chunksize(FD) )
    {
[3]     FD = FD->fd;
    }
[4] BK = FD->bk;
[5] P->bk = BK;
[6] P->fd = FD;
[7] FD->bk = BK->fd = P;
}
```

- Store chunk of size `S`, pointed to by `P` at appropriate position in the double linked list of bin with index `IDX`

- Macro

[1] `FD` initialized with a pointer to the start of the list of the given bin

[2] Loop searches the double linked list to find first chunk not larger than `P` or the end of the list by following consecutive forward pointers (line **3**)

[4] Follow back pointer `BK` to previous element in list

[5]+[6] Set backward and forward pointers for chunk `P`

[7] Update backward pointer of next chunk & forward pointer of previous chunk to address of chunk `P` (field `fd` at 8 byte offset within a boundary tag)

Unlink()

- Simplified version

```
#define unlink(P, BK, FD)
{
[1] FD = P->fd;
[2] BK = P->bk;
[3] FD->bk = BK;
[4] BK->fd = FD;
}
```

- Macro

- [1] [2] Save pointers in chunk **P** to **FD** and **BK**
- [3] Update backward pointer of next chunk in the list: address located at **FD** plus 12 bytes (offset of *bk* field in boundary tag) overwritten with value stored in **BK**
- [4] Update forward pointer of the previous chunk

4. Double-free vulnerabilities

4.1 Double Free Attack

- Call `free(A)` with forward consolidation after allocation of memory chunk **A** to create a larger chunk
- Allocate chunk **B**, hoping to get space just freed
- Copy ghost chunk into **B** at the location of **A** and a free ghost chunk adjacent to the chunk at **A**
 - ghost chunk is written into **B** and `free(A)` is called
- Calling `free(A)` again which coalesces the two ghost chunks and will try to remove the free ghost chunk from its bin

Removing ghost chunk

- Attacker writes fake forward and backward pointers into first 8 bytes of free ghost chunk
 - fd: target address to be overwritten, minus 12
 - bk: value to be written to the target address (fd + 12)
- Unlinking free ghost chunk uses fake pointers fd and bk
- Deallocation of heap is not atomic
 - Second call of `free(A)` only has the desired effect if the pointer to A had not been set to null when A was freed the first time
- Issue caused by function releasing memory but not setting pointer to null

4.2 Defences

- Stay true to abstraction: do not apply unlink to a chunk that is not part of a double link list
- Protect boundary tags with canaries
 - When chunk is allocated, initialise canary
 - When chunk is freed, recalculate checksum and compare with canary
- Randomise memory allocation so attacker cannot predict next allocated chunk
- Split user data from control data